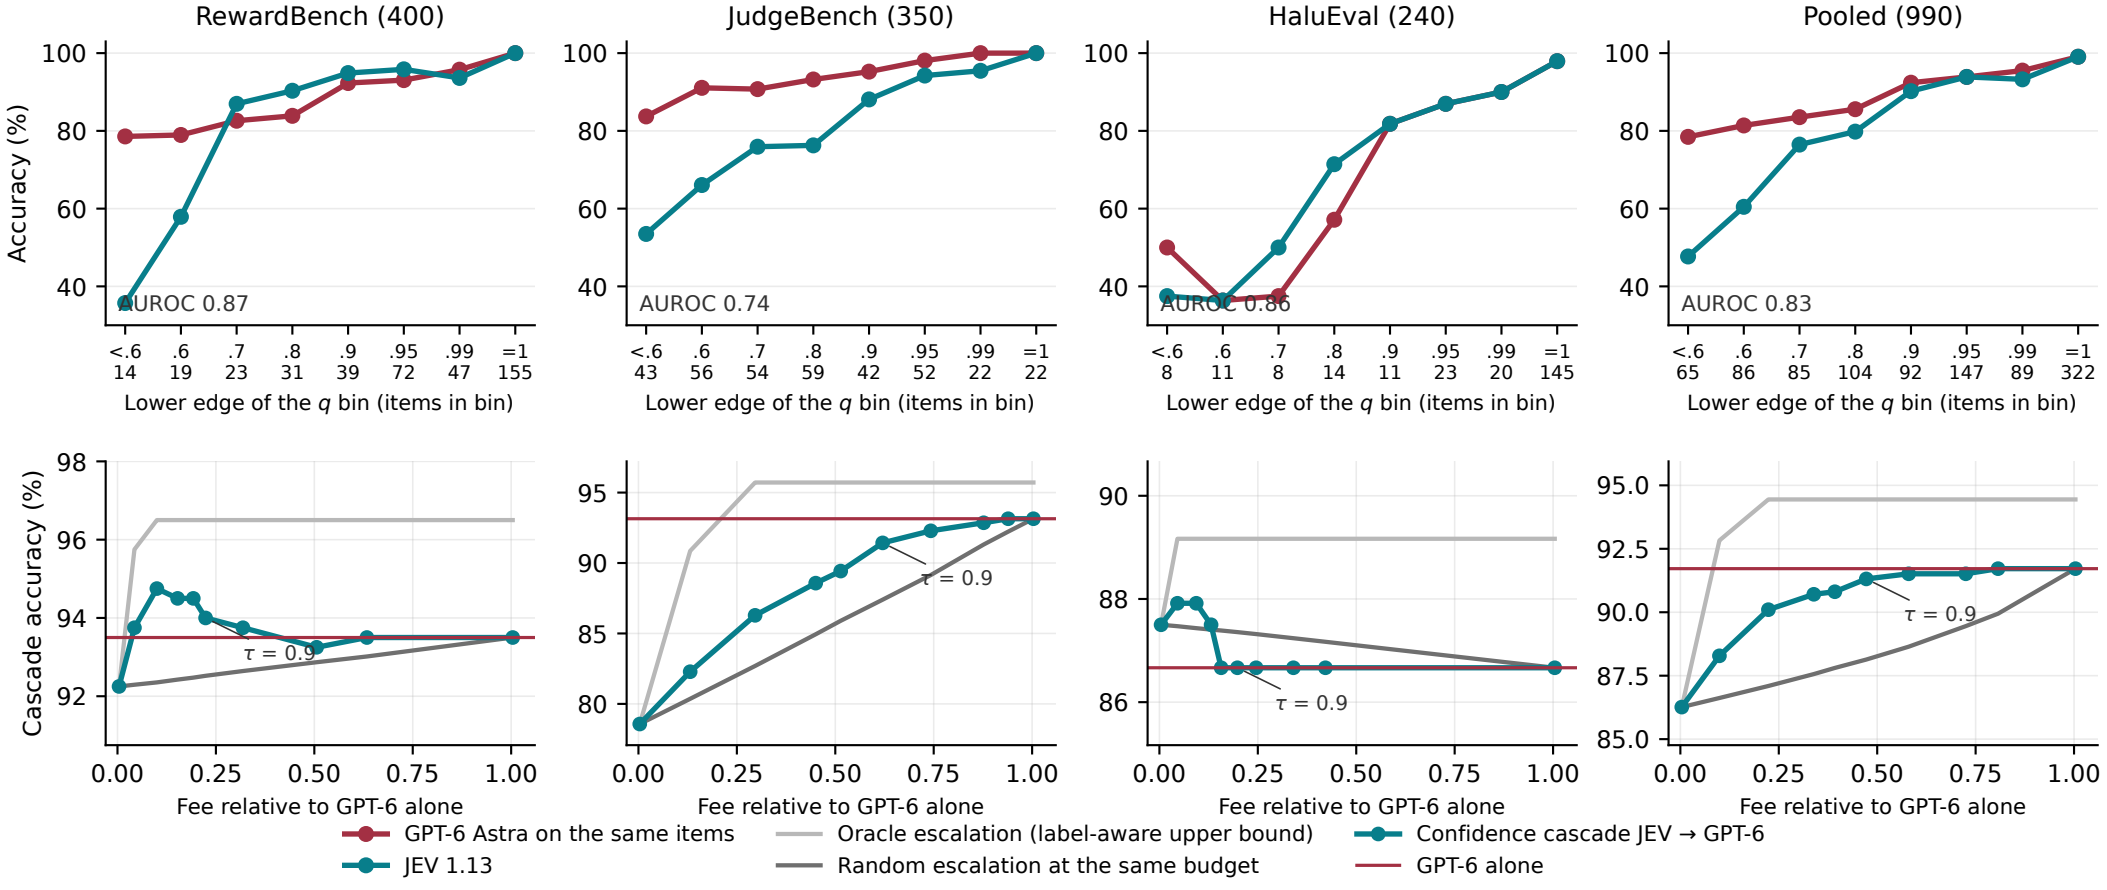
RewardBench (400)
JudgeBench (350)
HaluEval (240)
Pooled (990)
Accuracy (%)
AUROC 0.87
AUROC 0.74
AUROC 0.86
AUROC 0.83
<.6 .6 .7 .8 .9 .95 .99 =1
14 19 23 31 39 72 47 155
43 56 54 59 42 52 22 22
8 11 8 14 11 23 20 145
65 86 85 104 92 147 89 322
Lower edge of the q bin (items in bin)
Cascade accuracy (%)
τ = 0.9
Fee relative to GPT-6 alone
GPT-6 Astra on the same items
JEV 1.13
Oracle escalation (label-aware upper bound)
Random escalation at the same budget
Confidence cascade JEV → GPT-6
GPT-6 alone