

Thinking-length distributions with and without prefill





GPT-5.6-sol prefill (4 words)

Claude Opus 4.8 prefill (4 words)





Kimi K2.5 prefill (4 words)

Self-prefill control (4 words)





Kimi K3







Kimi K2.6







Kimi K2.5*







GLM-5.2



DeepSeek-V3.1



Inkling



thinking length (tokens, log scale)

control (no prefill) prefilled continuation prefix source